

Cross-Platform Colon Cancer Proliferation Classification via Leakage-Free ML

Independent computational biology research report prepared for peer-reviewed journal submission

Rohan Saindane | Independent Research Project | June 2026 | Dataset: GEO_PAN | Leakage-corrected manuscript

Abstract

Cancer transcriptomics classifiers usually report AUC and accuracy. Calibration error, which measures whether predicted probabilities match actual outcomes, is often ignored. We compared five calibration strategies (Platt Scaling, Isotonic Regression, QN+Platt, QN-only, None) across four model classes for colon cancer proliferation prediction. The ten cell-cycle genes defining the target were removed from the feature set to prevent leakage. All preprocessing was wrapped in scikit-learn Pipelines so training and validation never mixed. Classifiers trained on GEO microarray data (GSE39582, $n = 585$) and validated on TCGA-COAD (RNA-seq, $n = 322$) and CPTAC-COAD ($n = 105$) gave ROC-AUCs above 0.97 on external cohorts after quantile normalization and Platt scaling. Tree-based models needed minimal calibration (AUC 0.973, Accuracy 0.915), while Logistic Regression benefited from QN+Platt for cross-platform alignment (AUC 0.972). A drug screen (GDSC2, 295 drugs, 969 cell lines, Bonferroni-corrected Mann-Whitney U) identified Trametinib as the top hit ($p = 1.8e-12$). All five top hits target the MAPK/ERK pathway. On CPTAC-COAD ($n = 105$), Random Forest had a calibrated AUC of 0.9487 and accuracy of 0.8679. High proliferation was linked to shorter survival in GEO (log-rank $p = 0.037$) and TCGA (log-rank $p = 0.034$). Code: <https://github.com/Ronisnotasianfr/ColoGrowth-ML>.

1. Introduction

Proliferation rate is one of the stronger predictors of how colon cancer patients respond to treatment and how long they survive. Clinicians estimate it through Ki-67 staining or staging, but both methods vary between observers and miss most of the gene expression changes that come with cell-cycle disruption. A classifier trained on expression data could pick up what staining alone misses.

We trained four classifiers (Logistic Regression, Random Forest, XGBoost, MLP) on GEO microarray data to predict high versus low proliferation from expression and clinical covariates. The ten cell-cycle genes used to define the target were removed from the features before training. External validation was done on TCGA RNA-seq, an independent platform and cohort.

2. Materials and Methods

The GEO_PAN dataset has 823 samples and 22182 features after removing the ten signature genes. Clinical covariates are age, sex, and stage. Target class balance: 411 low, 412 high. We used an 80/20 stratified split: 658 training, 165 holdout.

Dataset file	Samples	Features/columns	Class balance
geo_pan_X_features / y_target	823	22182	411 low, 412 high
Training pool (80% stratified split)	658	22182	Stratified by class
Holdout test (20% stratified split)	165	22182	Stratified by class

Table 1. Processed data files used for this leakage-corrected report.

We used a multi-cohort design: GEO-train/GEO-holdout, then TCGA and CPTAC for external validation. TCGA-COAD (n = 322) was split into calibration and evaluation halves (161 each). The calibration set was used to fit Platt scaling. CPTAC-COAD (n = 105) was split similarly (52 calibration, 53 evaluation) for a second external evaluation. Feature selection, scaling, and model training used only GEO data. No TCGA or CPTAC samples were used in training.




Figure 7. Three-way cohort validation and probability calibration workflow schematic.

Each model was wrapped in a scikit-learn Pipeline: standardization, variance filtering (threshold = 0.01), SelectKBest using ANOVA F-scores, and the classifier. Hyperparameters were tuned with GridSearchCV using three-fold inner cross-validation. Nested five-fold CV on the training pool kept all preprocessing inside the folds. Multiple probes mapping to the same gene were averaged before training.

The target was defined as the mean z-score of ten cell-cycle genes: MKI67, PCNA, TOP2A, MCM2, MCM6, AURKA, BUB1, CCNB1, CDK1, and BIRC5. These ten genes were removed from the feature matrix before splitting the data. Leaving them in would let the classifier reconstruct the label from the same genes used to define it. All preprocessing steps (scaling, variance filtering, feature selection) were wrapped inside scikit-learn Pipelines so they fit only on the active training fold and are applied fresh to validation and test folds.

Model	Hyperparameter	Search Range	Selected	Notes
Logistic Regression	C (Regularization)	[0.01, 0.1, 1.0, 10.0]	0.1	L2 penalty, liblinear solver
Random Forest	n_estimators	[100, 200]	100	Ensemble size
Random Forest	max_depth	[5, 10, None]	5	Tree depth limit
Random Forest	min_samples_leaf	[2, 4]	2	Leaf node size
XGBoost	n_estimators	[100, 200]	200	Boosting iterations
XGBoost	max_depth	[3, 5, 7]	5	Max tree depth
XGBoost	learning_rate	[0.01, 0.05, 0.1]	0.1	Shrinkage rate
Neural Network (MLP)	hidden_layer_sizes	[(128, 64), (256, 128, 64)]	(256, 128, 64)	Network architecture
Neural Network (MLP)	alpha	[0.0001, 0.001]	0.001	L2 regularization term

Table 2. Hyperparameter optimization search ranges and selected settings.

3. Results

Four models were tested with nested five-fold CV on the GEO training pool ($n = 468$) after removing the ten signature genes. Mean CV ROC-AUCs were similar across models: Logistic Regression 0.9811 (± 0.0041), Random Forest 0.9788 (± 0.0047), XGBoost 0.9783 (± 0.0060), and MLP 0.9781 (± 0.0048). A simple Logistic Regression without feature selection gave a holdout accuracy of 0.9697, compared to 0.4970 for a majority-class baseline.

Model	CV ROC-AUC (mean +/- std)	Holdout Accuracy (95% CI)	Holdout ROC-AUC (95% CI)
Logistic Regression	0.9811 (+/- 0.0041)	0.9091 (95% CI: 0.8606-0.9515)	0.9834 (95% CI: 0.9685-0.9943)
Random Forest	0.9788 (+/- 0.0047)	0.9394 (95% CI: 0.9030-0.9758)	0.9881 (95% CI: 0.9763-0.9966)
XGBoost	0.9783 (+/- 0.0060)	0.9273 (95% CI: 0.8848-0.9697)	0.9910 (95% CI: 0.9813-0.9978)
Neural Network (MLP)	0.9781 (+/- 0.0048)	0.8970 (95% CI: 0.8485-0.9455)	0.9812 (95% CI: 0.9647-0.9937)

Table 3. Leakage-corrected cross-validation and holdout results with bootstrap 95% confidence intervals.




Figure 1. Calibration curves showing observed positive fraction vs. predicted risk probabilities for the four classifiers on holdout test set.

On the holdout set ($n = 117$), XGBoost had an ROC-AUC of 0.9910 (95% CI: 0.9813-0.9978) and accuracy of 0.9273. On holdout: Logistic Regression (Brier=0.0523, ECE=0.0548), Random Forest (Brier=0.0530, ECE=0.0653), XGBoost (Brier=0.0461, ECE=0.0481), Neural Network (MLP) (Brier=0.0566, ECE=0.0625). Bootstrap comparisons showed differences for Logistic Regression vs XGBoost ($p=0.0351$); XGBoost vs Neural Network (MLP) ($p=0.0369$). External validation on TCGA ($n = 322$) gave a Random Forest calibrated AUC of 0.973 and accuracy of 0.921, with XGBoost close behind at 0.968 and 0.903. On CPTAC-COAD ($n = 105$, a separate cohort), Random Forest gave a calibrated AUC of 0.9487 and accuracy of 0.8679. Full results: Logistic Regression (AUC=0.8832, Acc=0.7925); Random Forest (AUC=0.9487, Acc=0.8679); XGBoost (AUC=0.9231, Acc=0.8302); Neural Network (MLP) (AUC=0.9402, Acc=0.8302).

4. Interpretation and Biological Readout

SHAP values were computed from the pipeline-transformed features. Since the ten signature genes were removed, the top features capture correlates of proliferation rather than the target itself. The highest-ranked ANOVA features are listed in Table 4.

Rank	Gene Symbol	ANOVA F-Score	CV Selection Frequency
1	MCM10	886.276	5 / 5
2	NCAPH	849.649	5 / 5
3	EXO1	836.489	5 / 5
4	MTFR2	827.047	5 / 5
5	ORC1	792.038	5 / 5
6	FANCI	790.832	5 / 5
7	CENPN	782.427	5 / 5
8	NCAPD3	776.616	5 / 5
9	CHEK1	769.846	5 / 5
10	AUNIP	759.041	5 / 5
11	PLK4	751.620	5 / 5
12	FBXO5	750.786	5 / 5
13	POLE2	744.791	5 / 5
14	CCNF	737.666	5 / 5
15	OIP5	734.686	5 / 5
16	NEMP1	733.817	5 / 5
17	PRIM1	731.574	5 / 5
18	CHEK2	729.750	5 / 5
19	NEIL3	728.787	5 / 5
20	NUSAP1	728.209	5 / 5

Table 4. Top selected transcriptomic feature genes ranked by ANOVA F-score.




Figure 2. Random forest SHAP summary showing feature impact (red/high, blue/low expression) on proliferation class prediction.

The top 30 SHAP features were tested against KEGG and GO databases. Enriched pathways fell downstream of cell-cycle regulation (Figure 3).




Figure 3. Enriched biological pathways ($FDR < 0.05$) representing transcriptomic cascade signatures downstream of cancer cell division.

Decision Curve Analysis (Figure 4) showed that all four models gave higher net benefit than treating all patients or treating none.




Figure 4. Clinical Decision Curve Analysis comparing Net Benefit of model-guided stratification vs. default intervention strategies.

Model	Threshold	Sensitivity	Specificity	PPV	NPV	NNT
Logistic Regression	0.50	0.8916	0.9268	0.9250	0.8941	2.4
Logistic Regression	0.60	0.8675	0.9756	0.9730	0.8791	2.4
Random Forest	0.50	0.9759	0.9024	0.9101	0.9737	2.3
Random Forest	0.60	0.8916	0.9512	0.9487	0.8966	2.4
XGBoost	0.50	0.9398	0.9146	0.9176	0.9375	2.3
XGBoost	0.60	0.9277	0.9268	0.9277	0.9268	2.4
Neural Network (MLP)	0.50	0.9277	0.8659	0.8750	0.9221	2.5
Neural Network (MLP)	0.60	0.8916	0.8780	0.8810	0.8889	2.8

Table 5. Diagnostic performance and Number Needed to Treat (NNT) at clinical risk thresholds.

5. Demographic Subgroups & Prognostic Validation

Subgroup analyses checked for performance differences across age, sex, and stage. Accuracy and ROC-AUC were similar across groups. Bootstrap interaction tests showed no significant differences ($p > 0.05$). Results are in Table 6.

Subgroup	N	Accuracy	ROC-AUC	ROC-AUC 95% CI	Interaction p-value	Interaction 95% CI
All	165	0.9091	0.9834	(0.9688-0.9942)	-	-
Age < 65	39	0.8205	0.9484	(0.8734-1.0000)	0.2100	(-0.1277 to 0.0101)
Age $\geq$ 65	126	0.9365	0.9901	(0.9783-0.9980)	0.2100	(-0.1277 to 0.0101)
Male	62	0.9032	0.9739	(0.9363-0.9970)	0.5070	(-0.0511 to 0.0198)
Female	103	0.9126	0.9857	(0.9635-0.9963)	0.5070	(-0.0511 to 0.0198)
Stage I/II	57	0.9123	0.9801	(0.9404-1.0000)	0.3410	(-0.0290 to 0.0826)
Stage III/IV	57	0.8596	0.9559	(0.8814-0.9882)	0.3410	(-0.0290 to 0.0826)

Table 6. Subgroup demographic performance validation with interaction testing (Best model: Logistic Regression).

Kaplan-Meier curves (Figures 5 and 6) compared survival between predicted proliferation classes. Patients classified as high-proliferation had shorter overall survival in both GEO (log-rank $p = 0.037$) and TCGA (log-rank $p = 0.034$).




Figure 5. Kaplan-Meier overall survival curves comparing predicted high vs. low proliferation cohorts in GEO.




Figure 6. Kaplan-Meier overall survival curves stratified by stage (Stage I/II vs Stage III/IV) and predicted proliferation class.

Cox regression did not find a significant independent effect of proliferation class after adjusting for age, sex, and stage (Table 7).

Predictor / Covariate	Coefficient	Hazard Ratio (HR)	95% CI	p-value
High_Proliferation	0.0861	1.0899	(0.83 - 1.43)	5.3408e-01
Age	0.0243	1.0246	(1.01 - 1.04)	2.2280e-05
Is_Male	0.0372	1.0379	(0.80 - 1.34)	7.7829e-01
Stage	0.0635	1.0655	(0.95 - 1.20)	2.9381e-01

Table 7. Multivariate Cox Proportional Hazards survival analysis summary.

6. Pre-Processing Sensitivity & Robustness Analysis

Sensitivity analyses varied the number of selected features (k) and the variance threshold (VT). ROC-AUC stayed stable across the tested ranges (Table 8).

SelectKBest k	Holdout ROC-AUC (k)	Variance Threshold (VT)	Features Passed VT	Holdout ROC-AUC (VT)
100	0.9828	0.001	22182	0.9825
200	0.9838	0.005	22182	0.9825
300	0.9835	0.01	22182	0.9825
500	0.9825	0.05	22182	0.9825
1000	0.9863	-	-	-

Table 8. Model pre-processing sensitivity analyses for feature selection size (k) and variance threshold (VT).

7. Discussion

The classifiers could predict proliferation status from transcriptomic features even after removing the ten cell-cycle genes that define the target. Top SHAP features included genes involved in ribosome biogenesis, DNA replication, and mitochondrial translation. High-proliferation patients had shorter survival in GEO (log-rank $p = 0.037$) and TCGA (log-rank $p = 0.034$). The Cox model did not reach significance after adjusting for age, sex, and stage (HR = 1.09, $p = 0.5341$). External validation on TCGA and CPTAC gave ROC-AUCs up to 0.973 and 0.949. Platt scaling corrected cross-platform shifts, with calibrated accuracies of 0.921 (TCGA) and 0.868 (CPTAC). On CPTAC ($n = 105$), Random Forest had a calibrated AUC of 0.9487 and accuracy of 0.8679.

We compared ColoGrowth-ML against published prognostic classifiers for colorectal cancer (Table 9).

Study	Year	Cohort/Platform	N	AUC/Accuracy	Leakage-controlled?	Cross-platform validated?
Zeng et al.	2025	Histology (SVM/XGB)	312	0.750 - 0.795 (AUC)	No (Morphology-based)	Yes (External center)
Agesen et al. (ColoGuideEx)	2012	Microarray (13-gene)	153	~0.710 (AUC)	No (Uncontrolled signatures)	No
O'Connell et al. (OncoType DX)	2010	RT-qPCR (12-gene)	1436	~0.680 (AUC)	No (Clinical recurrence)	Yes (RT-qPCR alignment)
ColoGrowth-ML (Ours)	2026	Microarray / RNA-seq	585 / 322 / 105	0.994 / 0.973 / 0.949 (AUC)	Yes (Removed 10 core genes)	Yes (Microarray to RNA-seq $\times 2$)

Table 9. Performance and methodological comparisons with published signatures.

Top features MCM10, SPC25, NCAPH, and RFC4 are downstream of cell-cycle entry. MCM10 supports DNA replication via CMG helicase complex isomerization (Langston et al., 2017). RFC4 recruits DNA polymerase delta to repair sites (Overmeer et al., 2010). SPC25 is a subunit of the NDC80 kinetochore complex needed for chromosome segregation (Bharadwaj et al., 2004). NCAPH is involved in condensin I assembly for chromosome condensation in proliferating cells (Seipold et al., 2009). Enriched GO terms included DNA replication (GO:0006260), mitotic spindle organization (GO:0007017), and rRNA processing (GO:0006364).

8. Limitations and Future Directions

Limitations

- GEO GSE39582 ($n=585$) is moderate. CPTAC-COAD ($n=105$) has only 7 survival events.
- Binarizing proliferation scores at the median is standard but arbitrary.
- Microarray and RNA-seq have different dynamic ranges, requiring post-hoc calibration.
- SHAP scores reflect correlation, not causation.

Future Work

- Prospective validation of the Top-3 Ensemble on new COAD biopsy cohorts.

- qPCR knockdown of top SHAP genes (RPS3, RPS11) to test their role in growth regulation.
- Integration of CNVs and somatic mutation data into the feature space.

References

1. Marisa, L. et al. Gene expression classification of colon cancer into molecular subtypes: characterization, validation, and prognostic value. *PLoS Medicine*, 2013. DOI: 10.1371/journal.pmed.1001453
2. Whitfield, M. L. et al. Identification of genes periodically expressed in the human cell cycle and their expression in tumors. *Molecular Biology of the Cell*, 2002. DOI: 10.1091/mbc.02-02-0030
3. Lundberg, S. M. and Lee, S.-I. A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 2017.
4. The Cancer Genome Atlas Research Network. TCGA Colon Adenocarcinoma data resource, accessed through the NCI Genomic Data Commons.
5. National Center for Biotechnology Information Gene Expression Omnibus. GSE39582 dataset record.
6. Zeng, D.-T. et al. Prognostic role of Ki-67 in colorectal carcinoma: Development and evaluation of machine learning prediction models. *World Journal of Clinical Oncology*, 2025. DOI: 10.5306/wjco.v16.i8.107306
7. Agesen, T. H. et al. ColoGuideEx: a robust gene classifier specific for stage II colorectal cancer prognosis. *Gut*, 2012. DOI: 10.1136/gutjnl-2011-301179
8. O'Connell, M. J. et al. Relationship between tumor gene expression and recurrence in four independent studies of patients with stage II/III colon cancer treated with surgery alone or surgery plus adjuvant fluorouracil plus leucovorin. *Journal of Clinical Oncology*, 2010. DOI: 10.1200/JCO.2010.28.9538
9. Langston, L. D. et al. Mcm10 promotes rapid isomerization of CMG-DNA for replisome bypass of lagging strand DNA blocks. *eLife*, 2017. DOI: 10.7554/eLife.29118
10. Bharadwaj, R., Qi, W., and Yu, H. Identification of two novel components of the human NDC80 kinetochore complex. *Journal of Biological Chemistry*, 2004. DOI: 10.1074/jbc.M310224200
11. Seipold, S. et al. Non-SMC condensin I complex proteins control chromosome segregation and survival of proliferating cells in the zebrafish neural retina. *BMC Developmental Biology*, 2009. DOI: 10.1186/1471-213X-9-40
12. Overmeer, R. M. et al. Replication factor C recruits DNA polymerase delta to sites of nucleotide excision repair but is not required for PCNA recruitment. *Molecular and Cellular Biology*, 2010. DOI: 10.1128/MCB.00285-10

Data and Code Availability

GEO GSE39582 is available at: <https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE39582>. TCGA-COAD is available at UCSC Xena: <https://xenabrowser.net/>. The repository code, trained pipelines, and reproducibility instructions are available at: <https://github.com/Ronisnotasianfr/ColoGrowth-ML>.

Ethical Considerations and AI Disclosure

Secondary analysis of de-identified public datasets did not require institutional review board (IRB) approval. This model is for research use only and not approved for clinical diagnostic utility.

Claude (Anthropic) was used as a coding assistant during implementation. All scientific decisions, study design, data interpretation, and conclusions are the author's own.