

Machine Learning Prediction of Colon Cancer Proliferation Class from Gene Expression and Clinical Features

Independent computational biology research report prepared for external faculty peer review

Rohan Saindane | Independent Research Project | June 2026 | Dataset: GEO | Leakage-corrected draft

Abstract

Cellular proliferation rate is a fundamental hallmark of cancer and a critical determinant of tumor aggressiveness, clinical prognosis, and therapeutic response. In this study, we developed a rigorous, leakage-free machine learning framework to predict binary tumor proliferation class (high vs. low) from transcriptomic profiles and clinical covariates, utilizing the GEO microarray (GSE39582, $n = 585$) and TCGA-COAD RNA-seq ($n = 322$) cohorts. To prevent target leakage, we computed a baseline proliferation index from a established 10-gene cell-cycle signature (including MKI67, PCNA, and TOP2A) and strictly removed these genes from the feature space prior to model training. Our pipeline employs stratified 80/20 train/test splitting, nested 5-fold cross-validation with fold-local preprocessing (standardization, variance filtering, and ANOVA-based SelectKBest feature selection), and hyperparameter optimization. On the GEO cohort, Logistic Regression achieved the strongest holdout ROC-AUC of 0.9939 with a holdout accuracy of 0.9487. External cross-cohort validation (GEO trained, TCGA tested) demonstrated high discriminative generalization (ROC-AUC up to 0.9775) but highlighted a calibration shift (accuracy of ~ 0.520), highlighting a key challenge in cross-platform clinical translation. Finally, Kaplan-Meier survival analysis validated the clinical relevance of our computed proliferation groups, showing statistically significant overall survival differences in both GEO ($p = 0.037$) and TCGA ($p = 0.034$). These results demonstrate that downstream transcriptional cascades carry robust, generalizable signals reflecting cancer cell growth rates even when primary cell-cycle drivers are excluded.

1. Introduction

Cellular proliferation is a major biological hallmark of cancer progression and a primary indicator of tumor growth speed. In colon adenocarcinoma (COAD), assessing tumor cell division rates holds profound clinical value, directly correlating with patient survival outcomes, disease recurrence probability, and sensitivity to chemotherapeutic agents. In standard clinical practice, proliferation is estimated via histological staining techniques (e.g., Ki-67 immunohistochemistry) or staging systems. However, these manual methods can suffer from inter-observer variability and do not capture the broad, underlying transcriptomic changes associated with cell-cycle deregulation. A computational approach using machine learning applied to high-throughput gene expression datasets could provide an automated, objective method for tumor growth classification and reveal novel transcriptional correlates of aggressive tumor division.

This study develops a highly rigorous, leakage-free machine learning framework to classify colon cancer samples into high vs. low proliferation states using gene expression profiles (microarray and RNA-seq) and clinical covariates. A key challenge addressed is the prevention of target leakage: the 10 hallmark proliferation genes utilized to build the target classification metric were completely removed from the feature space prior to training. The pipeline compares Logistic Regression, Random Forest, XGBoost, and Multilayer Perceptron models, validating internally via nested cross-validation and externally across platforms (microarray to RNA-seq).

2. Materials and Methods

The GEO processed dataset comprises 585 samples and 22182 features after removing target-defining genes. Clinical covariates include age, sex, and tumor stage. The target label is balanced with 292 low, 293 high. The dataset was split into an 80% training pool (468 samples) and a 20% stratified holdout test split (117 samples).

Dataset file	Samples	Features/columns	Class balance
geo_X_features / y_target	585	22182	292 low, 293 high
Training pool (80% stratified split)	468	22182	Stratified by class
Holdout test (20% stratified split)	117	22182	Stratified by class

Table 1. Processed data files used for this leakage-corrected report.

To predict proliferation classes, four classifiers were wrapped in scikit-learn Pipelines to guarantee strict data separation during cross-validation. Standard scaling, low-variance feature filtering (threshold = 0.01), and SelectKBest feature selection (based on the ANOVA F-value) were fit exclusively on the training folds. Hyperparameters were optimized using GridSearchCV on the training pool.

A central challenge in clinical machine learning is target leakage, which occurs when information from the target variable is inadvertently included in the feature set. Here, the target class (high vs. low proliferation) was established using the mean z-score expression of 10 hallmark proliferation genes: MKI67, PCNA, TOP2A, MCM2, MCM6, AURKA, BUB1, CCNB1, CDK1, and BIRC5. If these genes were retained in the feature matrix, a classifier could trivially reconstruct the label with near-perfect accuracy, yielding scientifically meaningless results. To enforce rigor, all 10 signature genes were removed from the feature matrix before any train-test splitting. Furthermore, to prevent data leakage (where information from the validation fold spills into the training process), all preprocessing steps—including StandardScaler, VarianceThreshold, and SelectKBest feature selection—were encapsulated inside a unified scikit-learn Pipeline. This ensures that feature selection and scaling parameters are calculated solely on the active training folds during cross-validation and applied transitively to the validation/test folds.

3. Results

To establish internal model performance, we conducted a nested 5-fold cross-validation on the GEO microarray training pool ($n = 468$). Despite the strict exclusion of the 10 target-defining cell-cycle signature genes, all four machine learning architectures demonstrated exceptionally strong predictive capacity. The Random Forest classifier achieved a mean CV ROC-AUC of 0.9832 (± 0.0094), while XGBoost reached 0.9756 (± 0.0120). The linear baseline, Logistic Regression (L2-regularized), performed comparably with a mean CV ROC-AUC of 0.9801 (± 0.0092). The Neural Network (MLP) also exhibited high accuracy, reaching a mean CV ROC-AUC of 0.9711 (± 0.0184). This high performance within the GEO cohort suggests that cell proliferation induces widespread downstream transcriptomic changes, affecting hundreds of genes involved in DNA replication, translation, and cell-cycle progression, which are effectively captured by the models.

Model	CV ROC-AUC (mean +/- std)	Holdout Accuracy	Holdout ROC-AUC
Logistic Regression	0.9801 (+/- 0.0092)	0.9487	0.9939
Random Forest	0.9832 (+/- 0.0094)	0.9316	0.9845
XGBoost	0.9756 (+/- 0.0120)	0.9487	0.9915
Neural Network (MLP)	0.9711 (+/- 0.0184)	0.9316	0.9828

Table 2. Leakage-corrected cross-validation and holdout results.




Figure 1. Holdout ROC curves comparing the four trained classifiers.

On the independent holdout test split ($n = 117$), Logistic Regression achieved the highest holdout ROC-AUC of 0.9939 with a classification accuracy of 0.9487. When subjected to cross-cohort external validation—where the models trained on the GEO microarray were evaluated directly on the TCGA RNA-seq dataset—we observed a striking phenomenon. The Logistic Regression model retained an extremely high discriminative power with an external ROC-AUC of 0.9775, and XGBoost achieved an external ROC-AUC of 0.9071. However, the raw classification accuracy dropped to approximately 0.520 across models, and Brier scores hovered near 0.476. This discrepancy points to a severe calibration shift: the distribution of feature values and baseline expression levels differs dramatically between microarray and RNA-seq platforms. As a result, while the models maintain an excellent ability to rank samples from lowest to highest proliferation (high ROC-AUC), the default classification threshold of 0.5 is no longer aligned, resulting in many false classifications and highlight the necessity of post-hoc calibration in cross-platform clinical tools.

4. Interpretation and Biological Readout

To open the 'black box' of our machine learning models, we computed SHAP (SHapley Additive exPlanations) values for the pipeline-transformed features on the holdout split. These values reflect the marginal contribution of each gene feature to the model's final prediction score. Because the 10 direct cell-cycle signature genes were removed, the top SHAP features identify novel, indirect gene pathways associated with cancer cell proliferation rates.




Figure 2. Random forest SHAP summary from the leakage-corrected evaluation run.

To clinically validate our computationally derived proliferation classes, we conducted Kaplan-Meier overall survival analysis. Log-Rank tests were performed to compare the survival probabilities of the high vs. low proliferation cohorts.




Figure 3. Kaplan-Meier overall survival curves comparing predicted high vs. low proliferation cohorts (GEO cohort).

5. Discussion

By ensuring that feature selection and scaling are restricted to internal training folds, we avoided artificial inflation of model performance. The biological readouts and survival outcomes suggest that secondary transcriptional pathways can serve as strong surrogate markers for tumor growth rates.

This study demonstrates that colon cancer proliferation status can be predicted with high accuracy from transcriptomic features even after removing the core 10-gene cell-cycle signature. This finding suggests that proliferation is not an isolated cellular process but rather a driver of global transcriptional remodeling. The top features highlighted by SHAP analysis point to downstream transcriptional pathways, including ribosome biogenesis, mitochondrial translation, and metabolic enzymes, which support the energetic demands of rapidly dividing tumor cells. The clinical validity of our target labeling was further confirmed by survival analysis: patients categorized as high-proliferation had a statistically significant reduction in overall survival time compared to low-proliferation patients in both the GEO ($p = 0.037$) and TCGA ($p = 0.034$) cohorts. Crucially, the external validation results highlight that while rank-order risk prediction is highly generalizable across platforms (high AUC), absolute probability predictions are susceptible to cross-platform calibration shifts. Future translation of such classifiers into clinical diagnostics will require robust normalization techniques (such as combat batch-correction) or platform-specific threshold calibration.

6. Limitations and Next Steps

- Incorporate platform batch-correction algorithms (e.g., ComBat) to align microarray and RNA-seq feature distributions.
- Explore probability calibration techniques (e.g., Platt Scaling, Isotonic Regression) to improve cross-cohort accuracy.
- Integrate clinical features directly as model features to investigate potential synergistic effects on prediction.
- Perform wet-lab qPCR validation on top-performing SHAP gene targets.
- Conduct survival modeling (e.g., Cox Proportional Hazards) to evaluate proliferation as an independent prognostic factor.

References

1. Marisa, L. et al. Gene expression classification of colon cancer into molecular subtypes with distinct clinical and survival characteristics. PLoS Medicine, 2013.
2. Whitfield, M. L. et al. Identification of genes periodically expressed in the human cell cycle by microarray hybridization. Molecular Biology of the Cell, 2002.

3. Lundberg, S. M. and Lee, S.-I. A unified approach to interpreting model predictions. Advances in Neural Information Processing Systems, 2017.
4. The Cancer Genome Atlas Research Network. TCGA Colon Adenocarcinoma data resource, accessed through the NCI Genomic Data Commons.
5. National Center for Biotechnology Information Gene Expression Omnibus. GSE39582 dataset record.

Peer Review Contact Plan

This manuscript is prepared to request peer feedback from academic reviewers specializing in cancer genomics or biomedical machine learning. Reviewers will be asked to critique the leakage-free pipeline design, the calibration shift under cross-platform validation, and the physiological relevance of SHAP-selected genes.